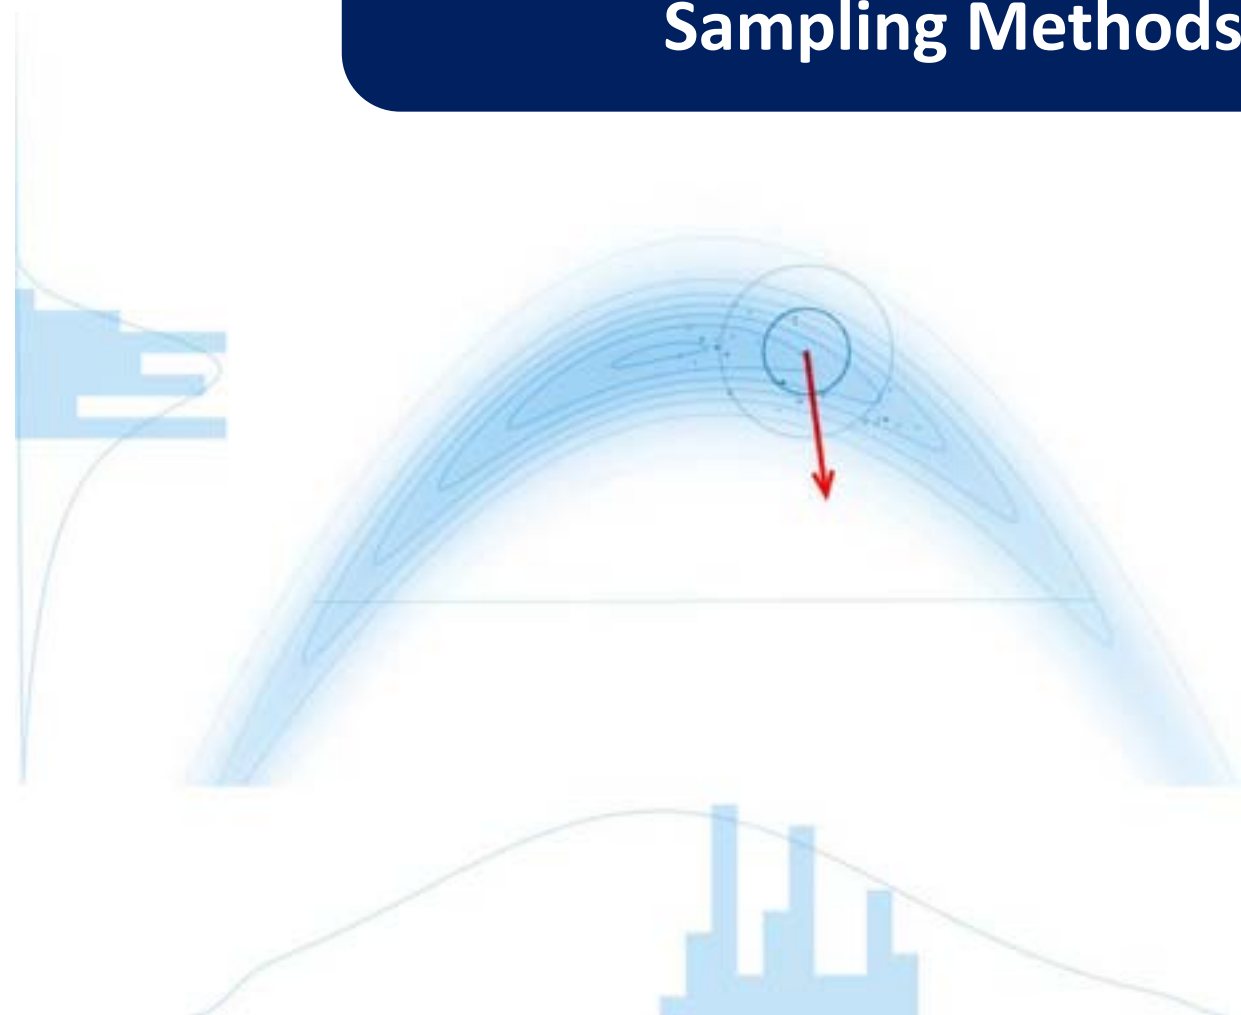


An Introduction to Monte Carlo Statistical Sampling Methods



Zhe Wei KHO
Jiacheng XU

The Bayesian Paradigm

The Posterior Function

$$p(\theta|\mathcal{D}) = \frac{p(\mathcal{D}|\theta)\pi(\theta)}{\int_{\Omega} p(\mathcal{D}|\theta)\pi(\theta) d\theta}$$

This is the starting point not the ending point for Bayesian modelling!

Posterior mean (estimate parameter values)

Credible Intervals (quantifying uncertainty)

Predictive posterior (predict new observations)

Bayesian Model Selection

Evaluating Utility Functions

All involve calculating
**integrals over the
posterior**

Integration is HARD!

What Do 'Real' Posterior Distributions Look Like?

Example: Statistical Model of Neuronal Assemblies

Likelihood

$$P(t, \omega, s \mid \vec{p}, \vec{\lambda}, \vec{n}) = \left(\prod_{i=1}^N n_{t_i} \right) \cdot \left(\prod_{\mu=1}^A \prod_{k=1}^M p_{\mu}^{\omega_{k\mu}} (1 - p_{\mu})^{1-\omega_{k\mu}} \right) \cdot \left(\prod_{i=1}^N \prod_{k=1}^M [\lambda_{t_i}(\omega_{kt_i})]^{s_{ik}} [1 - \lambda_{t_i}(\omega_{kt_i})]^{(1-s_{ik})} \right)$$

Priors

$$\begin{aligned} p_{\mu} &\sim \text{Beta}(\alpha_{\mu}^{(p)}, \beta_{\mu}^{(p)}) \\ \lambda_{\mu}(z) &\sim \text{Beta}(\alpha_{z,\mu}^{(\lambda)}, \beta_{z,\mu}^{(\lambda)}) \\ \{n_1, \dots, n_A\} &\sim \text{Dir}(\alpha_1^{(n)}, \dots, \alpha_A^{(n)}) \end{aligned}$$

Unlike textbook examples, Bayesian models that show up in research are typically **analytically intractable** and **high dimensional**!

The World of Monte Carlo



Throwing dice won't solve your financial woes...but maybe it can solve your mathematical ones?

Sampling

Suppose that we want to calculate the **mean value** of the **roll of a fair six-sided die**

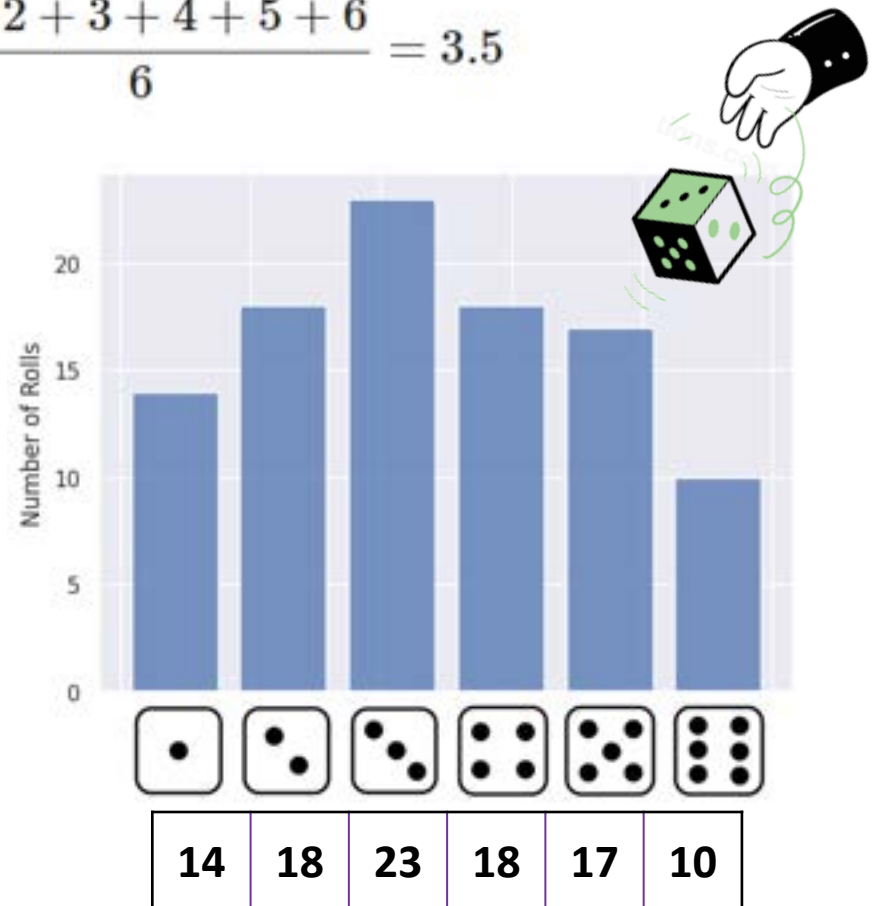
If we have **full information** about $p(x)$, we can simply compute

$$\mathbb{E}_p(X) = \sum_{x \in \mathcal{X}} xp(x) = \frac{1 + 2 + 3 + 4 + 5 + 6}{6} = 3.5$$


If we don't, then we can **estimate** it by **drawing samples from $p(x)$**

$$\bar{X} = \frac{1}{N} \sum_{i=1}^N x_i = 3.36$$

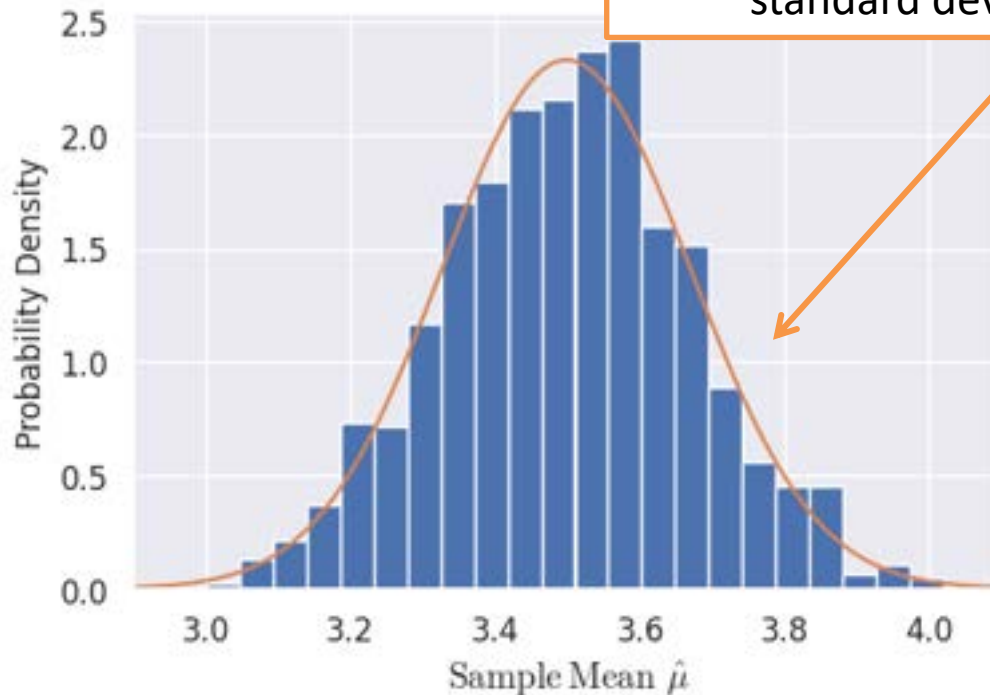
This is the **sample mean**



Sampling



Gaussian distribution with same mean and standard deviation as the data



Suppose we perform many sets of 100 throws.

The **sample mean** is a **random variable** as well.

The distribution of the sample mean approaches a normal distribution (**Central Limit Theorem**)

Law of the Unconscious Statistician

Law of the Unconscious Statistician

The expectation value of a function g of a random variable x is given by

$$\mathbb{E}[g(X)] = \sum_{x \in \mathcal{X}} g(x) p_X(x) \quad (\text{discrete})$$

$$\mathbb{E}[g(X)] = \int_{x \in \mathcal{X}} g(x) p_X(x) \mathrm{d}x \quad (\text{continuous})$$

(actually a nontrivial result but we will not prove it)

Strictly speaking, the expectation value of $g(X)$ should be taken wrt the PDF of $g(X)$.

We use this result pretty much every time we compute the posterior expectation of some function


Rejection Sampling

IDEA: As silly as it may sound, simply **prune the samples** until they fit the desired distribution!

Example: Selecting Samples Lying within the Unit Circle





"If it didn't hit the dartboard I didn't throw it. I didn't see anything and neither did you."




Rejection Sampling

We can imagine doing the same thing by treating the **target PDF** as a **target region**



Choosing a **proposal distribution** that more tightly bounds the target results in more efficient 'shooting'

How Do We Actually 'Shoot'?



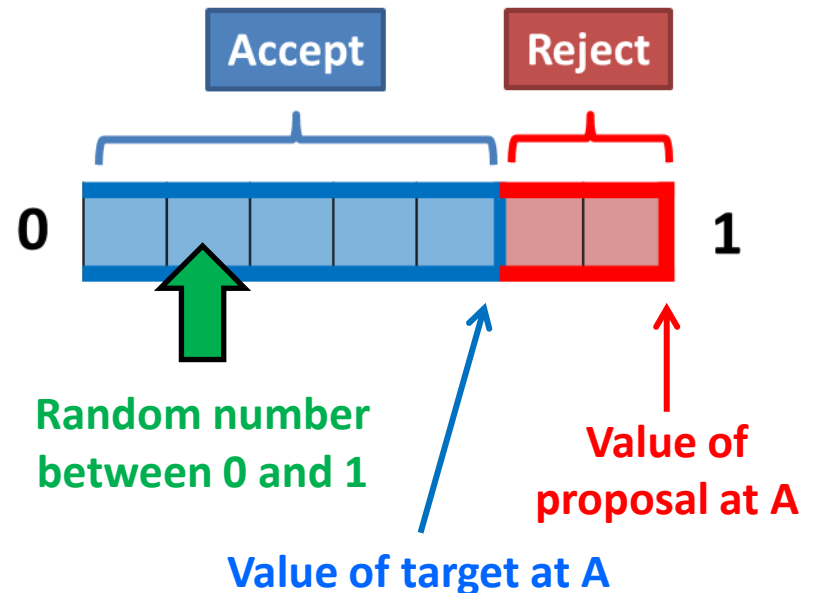
Height is proportionate to the **number of times** we land at that point

SOLUTION

Only keep:

$5/7$ of the samples at A

$3/7$ of the samples at B



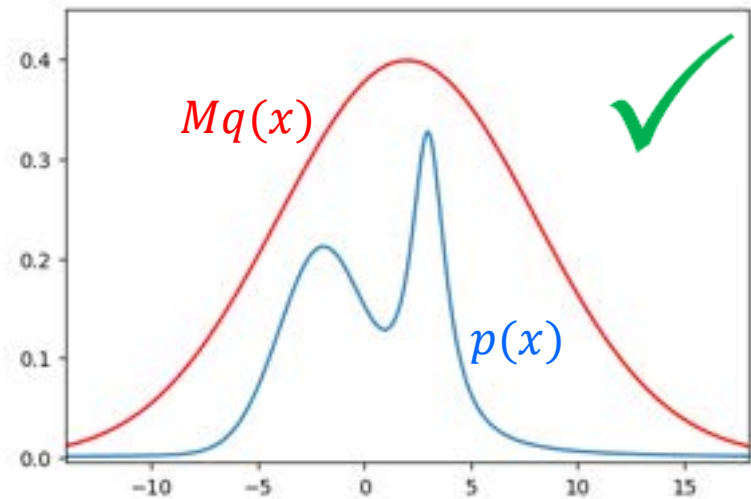

We **land** on A and B **equally often** but we **want** A and B to be in the **ratio 5:3**

We need to **reject** the **excess density**, and do so **probabilistically**

Rejection Sampling

Algorithm: Rejection Sampling

- ① Choose an appropriate proposal distribution q and scale factor M




Rescaled proposal distribution should be **larger** than the **target** – we can discard but not add!

Rejection Sampling

Algorithm: Rejection Sampling



- ② Simulate a candidate sample $x \sim q$ from the proposal density



Rejection Sampling

Algorithm: Rejection Sampling

- ① Choose an appropriate proposal distribution q and scale factor M
- ② Simulate a candidate sample $x \sim q$ from the proposal density





Rejection Sampling

Algorithm: Rejection Sampling

③

Calculate the **acceptance ratio** $r = p(x)/Mq(x)$





Rejection Sampling

Algorithm: Rejection Sampling

④

Simulate a random number $u \sim \mathcal{U}(0,1)$. If $u < r$, **accept** the sample. Otherwise, **reject** it.



Rejection Sampling







Algorithm: Rejection Sampling

- ① Choose an appropriate proposal distribution q and scale factor M
- ② Simulate a candidate sample $x \sim q$ from the proposal density
- ③ Calculate the acceptance ratio $r = p(x)/Mq(x)$
- ④ Simulate a random number $u \sim \mathcal{U}(0,1)$. If $u < r$, accept the sample. Otherwise, reject it.
- ⑤ Repeat until the desired number of samples are obtained.

Importance Sampling – A Dice Throwing Example



Regular Six-Sided Dice







						
# Rolls	14	18	23	18	17	10
Contribution	14	36	69	72	85	60

$$\mathbb{E}(X) = \frac{14 + 36 + 69 + 72 + 85 + 60}{14 + 18 + 23 + 18 + 17 + 10} = 3.36$$

Importance Sampling – A Dice Throwing Example



Regular Six-Sided Dice







						
# Rolls	14	18	23	18	17	10
Weights	1	1	1	1	1	1
Effective # Rolls	14	18	23	18	17	10
Contribution	14	36	69	72	85	60

How **many times** we **count each dice roll**! We will see what tricks we can play with this soon...

“Throwing” a Four-Sided Dice!



Four-Sided Dice

						
# Rolls	14	18	23	18	17	10
We				1	1	1
Effective # Rolls	14	18	23	18	17	10
Contribution	14	36	69	72	85	60







What if we simply don't count rolls that land on 5 or 6?!

How can we 'throw' a four-sided die when we only have a six-sided die?

“Throwing” a Four-Sided Dice!



Four-Sided Dice




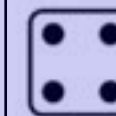


						
# Rolls	14	18	23	18	17	10
Weights	1	1	1	1	0	0
Effective # Rolls	14	18	23	18	0	0
Contribution	14	36	69	72	0	0

Change the weights to 0!

$$\mathbb{E}(X) = \frac{14 + 36 + 69 + 72 + 0 + 0}{14 + 18 + 23 + 18 + 0 + 0} = 2.62$$

“Throwing” a ‘Cheat’ Dice!



Cheat Dice						
						
# Rolls	14	18	23	18	17	10
Weights	1	1	1	1	1	1
Effective # Rolls	14	18	23	18	17	10
Contribution	14	36	69	72	85	60

Present 2 times

Not present

“Throwing” a ‘Cheat’ Dice!



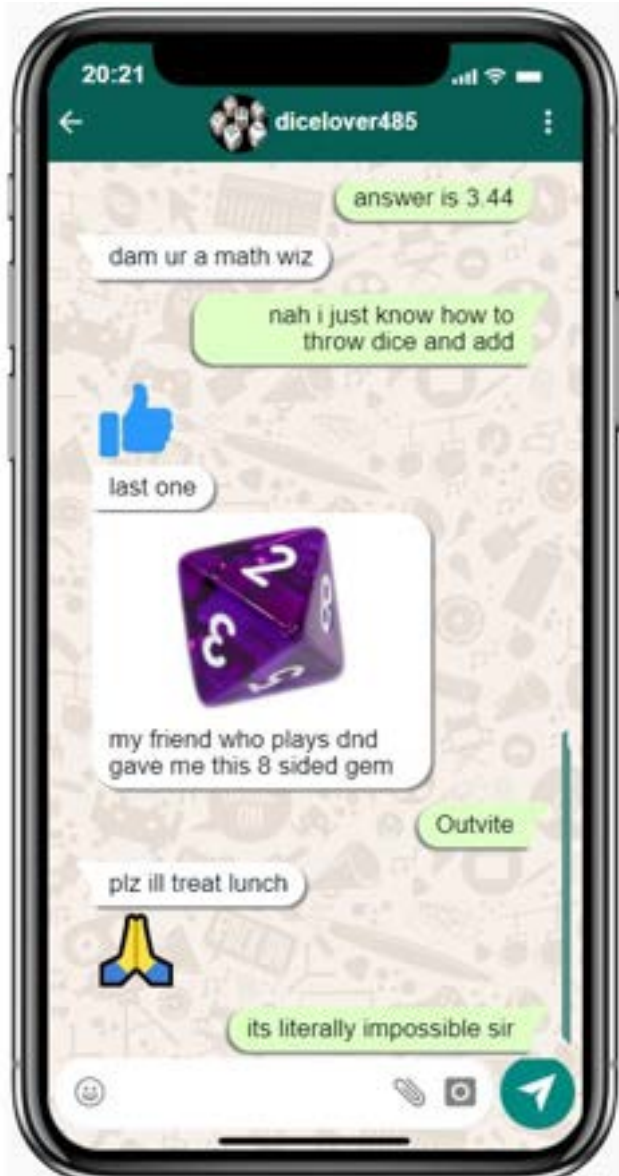
Cheat Dice

Double the weights!







# Rolls	14	18	23	18	17	10
Weights	0	1	2	2	0	1
Effective # Rolls	0	18	46	36	0	10
Contribution	0	36	138	144	0	60

$$\mathbb{E}(X) = \frac{0 + 36 + 138 + 144 + 0 + 60}{0 + 18 + 46 + 36 + 0 + 10} = 3.44$$

Eight-Sided Dice...?



Eight Sided Dice

							7	8
# Rolls	14	18	23	18	17	10	0	0
Weights	1	1	1	1	1	1	1?	1?
Effective # Rolls	14	18	23	18	17	10	0	0
Contribution	14	36	69	72	85	60	0	0

Impossible to get a sensible result if we **can never roll on them** to begin with!

Importance Sampling



Proposal Distribution $q(X)$	1/6	1/6	1/6	1/6	1/6	1/6
------------------------------	-----	-----	-----	-----	-----	-----



Target Distribution $p(X)$	1/4	1/4	1/4	1/4	0	0
Unnormalized Weights	3/2	3/2	3/2	3/2	0	0
Normalized Weights	1/4	1/4	1/4	1/4	0	0



Target Distribution $p(X)$	0	1/6	1/3	1/3	0	1/6
Unnormalized Weights	0	1	2	2	0	1
Normalized Weights	0	1/6	1/3	1/3	0	1/6

Unnormalized Weights $\tilde{w}(x) = \frac{p(x)}{q(x)}$

If you are sampling a state **too often**, count it **fewer times**!

Importance Sampling

NOTE:

Importance Sampling is **NOT** a sampling algorithm.
It is an **integration** algorithm

Conventional Monte Carlo Integration

$$\mathbb{E}_p[f(X)] = \int f(x) p(x) dx \approx \frac{1}{N} \sum_i^N f(x_i) \quad x_i \sim p(x)$$

Draw samples x_i from $p(x)$, calculate $p(x_i)$ and **average** them

Requires **being able to sample** from $p(x_i)$

If $p(x)$ is high where $f(x)$ is low, a lot of time is spent
sampling **unimportant points**

Importance Sampling

Importance Sampling

$$\mathbb{E}_p[f(X)] = \int f(x) p(x) dx = \int f(x) \frac{p(x)}{q(x)} q(x) dx$$

$$\approx \frac{1}{N} \sum_i^N f(x_i) \frac{p(x_i)}{q(x_i)}$$

$$= \frac{1}{N} \sum_i^N f(x_i) w(x_i)$$

Shifting **probability density** into **weights**

Distribution we know
to sample from


$$x_i \sim q(x)$$

Draw samples x_i from $q(x)$, calculate $p(x_i)$ and take a **weighted average**

Note: it is generally a good practice to **normalize** the weights
(even if they are normalized in principle)

This also implies that $p(x)$ does **NOT** need to be normalized

Importance Sampling



Height is proportionate to the **number of times** we land at that point

IDEA: If $q(x)$ samples a point too often compared to $p(x)$, simply “count” that point less often!

SOLUTION

We count:

Samples at A as **5/7** of a data point

Samples at B as **3/7** of a data point

We **land** on A and B **equally often** but we **want** A and B to be in the **ratio 5:3**


$$f(x_A) \rightarrow \frac{5}{7} f(x_A)$$

$$f(x_B) \rightarrow \frac{3}{7} f(x_B)$$

Importance Sampling

(HOMEWORK PROBLEM)

Example: Importance Sampling from the **Cauchy distribution** using a **Gaussian proposal**



Poor choice of importance function can lead to **exponentially increasing** weights!


If the weights are too **unbalanced**, then the algorithm is effectively only sampling $f(x)$ from points with high weights

Notion of **Effective Sample Size**:

$$n_e = \frac{(\sum_{i=1}^n w_i)^2}{\sum_{i=1}^n w_i^2}$$

Related to variance of the weights

Importance Sampling in D dimensions



Fraction of samples with **nonzero weights** decreases exponentially with **# of dimensions**

Most of our samples are **useless** for calculations!


Intuition about **geometry** in **high-dimensional space** is difficult

(Irrelevant) volume grows **exponentially** as **number of parameters** increases


This is called **The Curse of Dimensionality**

The Curse of Dimensionality

In **high dimensions**, most of the **volume** tends to be **concentrated** on the **edge of the sample space** (corners of the hypercube spanned by the variables)




No. of Dimensions	V_{sphere} / V_{cube}
3	0.52360
4	0.30843
5	0.16450
...	...
10	0.00249
...	...
20	2.4611×10^{-8}



The Typical Set – Where All The Probability Is

$$\mathbb{E}_p[f(X)] = \int f(x) \underline{p(x)} dx$$



Only the **typical set** has non-negligible contributions

Want to focus our sampling efforts there, but **HOW?**

The Markov Chain Monte Carlo Revolution

We have this ugly posterior that has inverse Gamma, Pareto, hypergeometric and god knows that other distributions in it. And it has 20 parameters. Any ideas what to do with it so we can get this damn paper published?



What if we construct an **ergodic, reversible Markov chain** that has an **equilibrium distribution which matches our posterior distribution?**

Make a simpler model

Apply for HPC time




The Blessing of Dimensionality?



Typical set is a **low(er) dimension manifold!**
(Preview: Gibbs sampling makes use of this fact!)

GOAL: Design something that **moves towards** the **typical set** and **stays there**, exploring it **thoroughly**

Random walk in the **state space** of the **typical set**




1 → 2 → 5 → 8 →
9 → 6 → 5 → ...

Markov Chains

Markov Chain

A system that experiences **transitions** from one state to another **probabilistically** & has **no memory**




Rows sum to one

$$\mathcal{T} = \begin{bmatrix} \mathbb{P}(1 \rightarrow 1) & \mathbb{P}(1 \rightarrow 2) & \mathbb{P}(1 \rightarrow 3) \\ \mathbb{P}(2 \rightarrow 1) & \mathbb{P}(2 \rightarrow 2) & \mathbb{P}(2 \rightarrow 3) \\ \mathbb{P}(3 \rightarrow 1) & \mathbb{P}(3 \rightarrow 2) & \mathbb{P}(3 \rightarrow 3) \end{bmatrix} = \begin{bmatrix} 0 & 1/2 & 1/2 \\ 1/4 & 1/2 & 1/4 \\ 1/4 & 1/4 & 1/2 \end{bmatrix}$$

Transition Matrix (Kernel for continuous case)

The Stationary Distribution of a Markov Chain



Sequence of States:

$1 \rightarrow 2 \rightarrow 3 \rightarrow 3 \rightarrow 2 \rightarrow 3 \rightarrow$
 $\dots \rightarrow 1 \rightarrow 1 \rightarrow 2 \rightarrow 2 \rightarrow 3 \rightarrow 2$

If we run the chain **long enough**, we will find that:

The system still **transitions** between the three states but the **asymptotic distribution** does not change


This is the **stationary state** of the system

$$\mathbb{P}(X = 1) \rightarrow \frac{1}{5}$$

$$\mathbb{P}(X = 2) \rightarrow \frac{2}{5}$$

$$\mathbb{P}(X = 3) \rightarrow \frac{2}{5}$$

Detailed Balance & No Probability Flow



$$\mathbb{P}(X = 1) \mathcal{T}(2 | 1) = \frac{1}{5} \cdot \frac{1}{2} = \frac{1}{10}$$

$$\mathbb{P}(X = 2) \mathcal{T}(1 | 2) = \frac{2}{5} \cdot \frac{1}{4} = \frac{1}{10}$$


Probability flow from $1 \rightarrow 2$ is the **same** as that from $2 \rightarrow 1$!

Detailed Balance

$$p(\theta) \mathcal{T}(\theta^* | \theta) = p(\theta^*) \mathcal{T}(\theta | \theta^*)$$

IDEA: Play around with the **transition matrix** to engineer the **distribution** we want!

Engineering Distributions



PROBLEM:

Jerry wants to spend time at the three locations **proportional to the amount of cheese** present.


However, he has no **sense of direction** and **cannot remember**.

The only thing he knows is **counting the number of cheese present**

TARGET:

$$P(A) = \frac{1}{3} \quad P(B) = \frac{1}{2} \quad P(C) = \frac{1}{6}$$

Engineering Distributions



If Jerry simply wanders around **randomly**, he would **spend equal time** at each location

IDEA:


If Jerry **compares** the amount of cheese at his **current location** and the **new location**, can he use this information to do anything?

Recall **Detailed Balance**

$$\frac{P(A)}{P(C)} = \frac{\mathcal{T}(C|A)}{\mathcal{T}(A|C)} = 2$$

Example: Jerry wants to somehow move from A to C **half as often** as moving from C to A

Engineering Distributions



Earlier, we saw how to adjust distributions to fit what we want. In this case, the base distribution is that of the **random walk** (which is uniformly distributed)

SOLUTION

If we have twice as many samples at C as what we wanted, simply **reject** half the moves from A to C (and keep all the moves from C to A)!!

If the new location has **more cheese**, Jerry **always moves** to it.

If the new location has **fewer cheese**, Jerry moves to it with **probability r** where r is the **ratio of the cheese at the new location to the old location**.

Engineering Distributions

Cheese Ratios

$$P(A) \propto 2 \quad P(B) \propto 1 \quad P(C) \propto 3$$

$$\mathcal{T} = \begin{bmatrix} \mathbb{P}(A \rightarrow A) & \mathbb{P}(A \rightarrow B) & \mathbb{P}(A \rightarrow C) \\ \mathbb{P}(B \rightarrow A) & \mathbb{P}(B \rightarrow B) & \mathbb{P}(B \rightarrow C) \\ \mathbb{P}(C \rightarrow A) & \mathbb{P}(C \rightarrow B) & \mathbb{P}(C \rightarrow C) \end{bmatrix} = \begin{bmatrix} 0 & 1/2 & 1/2 \\ 1/2 & 0 & 1/2 \\ 1/2 & 1/2 & 0 \end{bmatrix}$$



Original Random Walk Matrix

Engineering Distributions

Cheese Ratios

$$P(A) \propto 2 \quad P(B) \propto 1 \quad P(C) \propto 3$$

$$\mathcal{T} = \begin{bmatrix} \mathbb{P}(A \rightarrow A) & \mathbb{P}(A \rightarrow B) & \mathbb{P}(A \rightarrow C) \\ \mathbb{P}(B \rightarrow A) & \mathbb{P}(B \rightarrow B) & \mathbb{P}(B \rightarrow C) \\ \mathbb{P}(C \rightarrow A) & \mathbb{P}(C \rightarrow B) & \mathbb{P}(C \rightarrow C) \end{bmatrix} = \begin{bmatrix} 0 & 1/2 & 1/2 \\ 1/2 & 0 & 1/2 \\ 1/2 & 1/2 & 0 \end{bmatrix}$$

Move to B half the time and **stay at A** (move back to A) the **other half of the time**

Engineering Distributions

Cheese Ratios

$$P(A) \propto 2 \quad P(B) \propto 1 \quad P(C) \propto 3$$

$$\mathcal{T} = \begin{bmatrix} \mathbb{P}(A \rightarrow A) & \mathbb{P}(A \rightarrow B) & \mathbb{P}(A \rightarrow C) \\ \mathbb{P}(B \rightarrow A) & \mathbb{P}(B \rightarrow B) & \mathbb{P}(B \rightarrow C) \\ \mathbb{P}(C \rightarrow A) & \mathbb{P}(C \rightarrow B) & \mathbb{P}(C \rightarrow C) \end{bmatrix} = \begin{bmatrix} 1/4 & 1/4 & 1/2 \\ 1/2 & 0 & 1/2 \\ 1/2 & 1/2 & 0 \end{bmatrix}$$

Move to B half the time and **stay at A** (move back to A) the **other half of the time**

Engineering Distributions

Cheese Ratios

$$P(A) \propto 2 \quad P(B) \propto 1 \quad P(C) \propto 3$$

$$\mathcal{T} = \begin{bmatrix} \mathbb{P}(A \rightarrow A) & \mathbb{P}(A \rightarrow B) & \mathbb{P}(A \rightarrow C) \\ \mathbb{P}(B \rightarrow A) & \mathbb{P}(B \rightarrow B) & \mathbb{P}(B \rightarrow C) \\ \mathbb{P}(C \rightarrow A) & \mathbb{P}(C \rightarrow B) & \mathbb{P}(C \rightarrow C) \end{bmatrix} = \begin{bmatrix} 1/4 & 1/4 & 1/2 \\ 1/2 & 0 & 1/2 \\ 1/2 & 1/2 & 0 \end{bmatrix}$$

Moves are **always accepted**

Engineering Distributions

Cheese Ratios

$$P(A) \propto 2 \quad P(B) \propto 1 \quad P(C) \propto 3$$

$$\mathcal{T} = \begin{bmatrix} \mathbb{P}(A \rightarrow A) & \mathbb{P}(A \rightarrow B) & \mathbb{P}(A \rightarrow C) \\ \mathbb{P}(B \rightarrow A) & \mathbb{P}(B \rightarrow B) & \mathbb{P}(B \rightarrow C) \\ \mathbb{P}(C \rightarrow A) & \mathbb{P}(C \rightarrow B) & \mathbb{P}(C \rightarrow C) \end{bmatrix} = \begin{bmatrix} 1/4 & 1/4 & 1/2 \\ 1/2 & 0 & 1/2 \\ 1/2 & 1/2 & 0 \end{bmatrix}$$

Move to A 2/3 the time and **stay at C** the other 1/3 of the time

Engineering Distributions

Cheese Ratios

$$P(A) \propto 2 \quad P(B) \propto 1 \quad P(C) \propto 3$$

$$\mathcal{T} = \begin{bmatrix} \mathbb{P}(A \rightarrow A) & \mathbb{P}(A \rightarrow B) & \mathbb{P}(A \rightarrow C) \\ \mathbb{P}(B \rightarrow A) & \mathbb{P}(B \rightarrow B) & \mathbb{P}(B \rightarrow C) \\ \mathbb{P}(C \rightarrow A) & \mathbb{P}(C \rightarrow B) & \mathbb{P}(C \rightarrow C) \end{bmatrix} = \begin{bmatrix} 1/4 & 1/4 & 1/2 \\ 1/2 & 0 & 1/2 \\ \mathbf{1/3} & 1/2 & \mathbf{1/6} \end{bmatrix}$$

Move to A 2/3 the time and **stay at C** the other 1/3 of the time

Engineering Distributions

Cheese Ratios

$$P(A) \propto 2 \quad P(B) \propto 1 \quad P(C) \propto 3$$

$$\mathcal{T} = \begin{bmatrix} \mathbb{P}(A \rightarrow A) & \mathbb{P}(A \rightarrow B) & \mathbb{P}(A \rightarrow C) \\ \mathbb{P}(B \rightarrow A) & \mathbb{P}(B \rightarrow B) & \mathbb{P}(B \rightarrow C) \\ \mathbb{P}(C \rightarrow A) & \mathbb{P}(C \rightarrow B) & \mathbb{P}(C \rightarrow C) \end{bmatrix} = \begin{bmatrix} 1/4 & 1/4 & 1/2 \\ 1/2 & 0 & 1/2 \\ 1/3 & 1/2 & 1/6 \end{bmatrix}$$

Move to B 1/3 the time and **stay at C** the other 2/3 of the time

Engineering Distributions

Cheese Ratios

$$P(A) \propto 2 \quad P(B) \propto 1 \quad P(C) \propto 3$$

$$\mathcal{T} = \begin{bmatrix} \mathbb{P}(A \rightarrow A) & \mathbb{P}(A \rightarrow B) & \mathbb{P}(A \rightarrow C) \\ \mathbb{P}(B \rightarrow A) & \mathbb{P}(B \rightarrow B) & \mathbb{P}(B \rightarrow C) \\ \mathbb{P}(C \rightarrow A) & \mathbb{P}(C \rightarrow B) & \mathbb{P}(C \rightarrow C) \end{bmatrix} = \begin{bmatrix} 1/4 & 1/4 & 1/2 \\ 1/2 & 0 & 1/2 \\ 1/3 & 1/6 & 1/2 \end{bmatrix}$$

Move to B 1/3 the time and **stay at C** the other 2/3 of the time

Engineering Distributions

Proposal Kernel

$$\begin{bmatrix} 0 & 1/2 & 1/2 \\ 1/2 & 0 & 1/2 \\ 1/2 & 1/2 & 0 \end{bmatrix}$$

Accept-Reject based
on the ratio of cheese
(**Metropolis ratio**)

Target Kernel

$$\begin{bmatrix} 1/4 & 1/4 & 1/2 \\ 1/2 & 0 & 1/2 \\ 1/3 & 1/6 & 1/2 \end{bmatrix}$$

$$q(\theta^* | \theta) \alpha(\theta^* | \theta) = \mathcal{T}(\theta^* | \theta)$$

**Scheme for proposing
new moves**

Final transition probability
(that has the target distribution
as its stationary distribution!)

$$\min [1, p(\theta)/p(\theta^*)]$$

This is where the
magic happens

**Accept if new move is 'better', otherwise
accept with probability r**

Metropolis Algorithm


Algorithm: Metropolis MCMC

- ① Choose an appropriate transition kernel $q(\theta^*|\theta)$

This determines how **new moves** are **proposed**

A common choice is the Gaussian **centered around the current location** (**Random Walk** Metropolis algorithm)

The variance of the Gaussian should be carefully chosen



Note: in vanilla Metropolis, this is a **symmetric function**. Metropolis-Hastings generalizes this to an arbitrary kernel.

Metropolis Algorithm

Algorithm: Metropolis MCMC

- ② Choose an appropriate initialization of the parameters

This is the **start location** of the algorithm

A poor choice far away from the typical set can lead to slow convergence.

A good guess is near the mode of the target distribution
(e.g. MAP of the posterior)

Metropolis Algorithm

Algorithm: Metropolis MCMC

- ③ Draw a sample θ^* from $q(\theta^*|\theta)$ and propose it as the next move

It goes without saying that $q(\theta^*|\theta)$ should be something you know how to sample from

Metropolis Algorithm

Algorithm: Metropolis MCMC

④

Calculate the Metropolis ratio $r = p(\theta)/p(\theta^*)$

This is basically comparing the relative probability densities of the current and proposed points

Metropolis Algorithm

Algorithm: Metropolis MCMC

- ⑤ Accept the proposed move with probability $\alpha(\theta^*|\theta) = \min[1, r]$

Always move to regions with **higher probability**

Sometimes move to regions with **lower probability**

IMPORTANT: Not moving means **sampling the current state again**

The acceptance probability should not be too high or too low.

Too high: proposed moves are likely to be nearby and system doesn't explore the sample space

Too low: system stays in same state for a long time and makes abrupt jumps

Metropolis Algorithm

Algorithm: Metropolis MCMC

- ① Choose an appropriate transition kernel $q(\theta^*|\theta)$
- ② Choose an appropriate initialization of the parameters
- ③ Draw a sample θ^* from $q(\theta^*|\theta)$ and propose it as the next move
- ④ Calculate the Metropolis ratio $r = p(\theta)/p(\theta^*)$
- ⑤ Accept the proposed move with probability $\alpha(\theta^*|\theta) = \min[1, r]$
- ⑥ Repeat until the desired number of samples are obtained.

Two algorithms

To get posterior,
normalization is hard:

$$p(\theta_1, \theta_2 \dots | D) = p(D | \theta_1, \theta_2 \dots) p(\theta_1, \theta_2 \dots) / \underline{p(D)}$$

Ways around: 1. **Metropolis:**
compare two points
(candidate vs. current)

$$\frac{p(D | \theta_1, \theta_2 \dots) p(\theta_1, \theta_2 \dots) / \cancel{p(D)}}{p(D | \theta_1^*, \theta_2^* \dots) p(\theta_1^*, \theta_2^* \dots) / \cancel{p(D)}}$$

2. **Gibbs:** deal with
one para at a time
 $p(\theta_i | \{\theta_{j \neq i}\}, D)$

To avoid the curse of
dimensionality

Compare by an example

Two coins example

- Estimate the biases θ_1, θ_2 of two independent coins

Coin #1



Coin #2



- Observations D:

$$\frac{\text{The total \# of heads}}{\text{The total flips}} \rightarrow \frac{z_1}{N_1} \quad \frac{z_2}{N_2}$$


- How to estimate θ_1, θ_2 ?

Bayesian:


$$\underline{p(\theta_1, \theta_2 | D)} = \underline{p(D | \theta_1, \theta_2)} \underline{p(\theta_1, \theta_2)} / \underline{p(D)}$$

Known up to a scale


Prior



Likelihood



Posterior



Two coins example

- Estimate the biases θ_1, θ_2 of two independent coins

Coin #1



Coin #2



- Observations D:

$$\frac{\text{The total \# of heads}}{\text{The total flips}} \rightarrow \frac{z_1}{N_1} \quad \frac{z_2}{N_2}$$


- How to estimate θ_1, θ_2 ?

Bayesian:


$$\underline{p(\theta_1, \theta_2 | D)} = \underline{p(D | \theta_1, \theta_2)} \underline{p(\theta_1, \theta_2)} / \underline{p(D)}$$

Known up to a scale


Prior



Likelihood



Posterior




Two coins example


Metropolis

1. Random initialization: θ_1, θ_2
2. Draw a candidate from the proposal distribution (a bivariate normal)
3. Get acceptance rate
4. Accept or reject
5. Repeat step 2

Known up to a scale
Posterior



Eff.Sz. θ_1 = 276, Eff.Sz. θ_2 = 253.1




Low efficiency, why?

Two coins example


Metropolis

1. Random initialization: θ_1, θ_2
2. Draw a candidate from the proposal distribution (a bivariate normal)
3. Get acceptance rate
4. Accept or reject
5. Repeat step 2

Known up to a scale
Posterior







Eff.Sz. θ_1 = 276, Eff.Sz. θ_2 = 253.1



Low efficiency, why?


Two coins example

Metropolis - limitations



Two coins example


Metropolis - limitations



Two coins example


Gibbs

1. Random initialization: θ_1, θ_2
2. **Fix** θ_2
3. Generate an **updated** θ_1 from the proposal distribution $p(\theta_1 | \theta_2, D)$ (Always accept)




How to generate θ_1 from $p(\theta_1 | \theta_2, D)$?

Rejection sampling



Importance sampling





Other
methods

Two coins example

Gibbs


1. Random initialization: θ_1, θ_2
2. **Fix** θ_2
3. Generate an **updated** θ_1 from the proposal distribution $p(\theta_1 | \theta_2, D)$ (Always accept)
4. **Fix** θ_1
5. Generate an **updated** θ_2 from the proposal distribution $p(\theta_2 | \theta_1, D)$ (Always accept)
6. Repeat step 2.



Two coins example

Gibbs

1. Random initialization: θ_1, θ_2
2. **Fix** θ_2
3. Generate an **updated** θ_1 from the proposal distribution $p(\theta_1 | \theta_2, D)$ (Always accept)
4. **Fix** θ_1
5. Generate an **updated** θ_2 from the proposal distribution $p(\theta_2 | \theta_1, D)$ (Always accept)
6. Repeat step 2.



Gibbs

- Random initialization
- Cycle through each of $\theta_1, \theta_2, \theta_3, \theta_4, \dots, \theta_1, \theta_2, \theta_3, \theta_4, \dots$
fix the rest
- Draw a candidate from the **proposal distribution**
 $p(\theta_i | \{\theta_{j \neq i}\}, D)$
- **Always accepted**
- Repeat

Metropolis

- Random initialization
- Draw a candidate from the **proposal distribution (e.g. a N-dimension normal)**
- **Get acceptance rate**
- **Determine to accept or not**
- Repeat

More efficient in high dimensions

Limit

- Need conditional probabilities
- Bad for highly correlated parameters

- **Tune step size.** (Proposal distribution similar to posterior distribution)
- **rejection high**
- **Trapped in local optimal**

Two coins example - final question

- Observations D:

Coin #1




$$\frac{z_1}{N_1}$$

Coin #2




$$\frac{z_2}{N_2}$$

- Estimate the independent biases, θ_1 , θ_2 by bayesian and MCMC
- Do two biases differ?**







For each pair of data, get the difference



How to evaluate the performance?


Performance evaluation

Sample more gives “Accuracy”




Performance evaluation




Sample more gives “Accuracy”



To quantify correlation




Autocorrelation ACF(k): how similar a trace is to itself with a time lag of k.



$$ACF(k) = \frac{\sum (x_t - \bar{x})(x_{t+k} - \bar{x})}{\sum (x_t - \bar{x})^2}$$

Performance evaluation

Sample more gives “Accuracy”



Lower autocorrelation,
better sampling


Quantify:
effective sample size

$$ESS = N / \left(1 + 2 \sum_{k=1}^{\infty} ACF(k) \right)$$

Denominator: Total ACF
from $-\infty$ and ∞

Performance evaluation

Sample more gives “Accuracy”



Intuitively, effective sample size (ESS) $\rightarrow \infty$,
perfect estimation

Quantify: **Monte Carlo standard error (MCSE)**


$$\text{MCSE} = \text{SD} / \sqrt{\text{ESS}}$$

$$\text{SE} = \text{SD} / \sqrt{N}$$

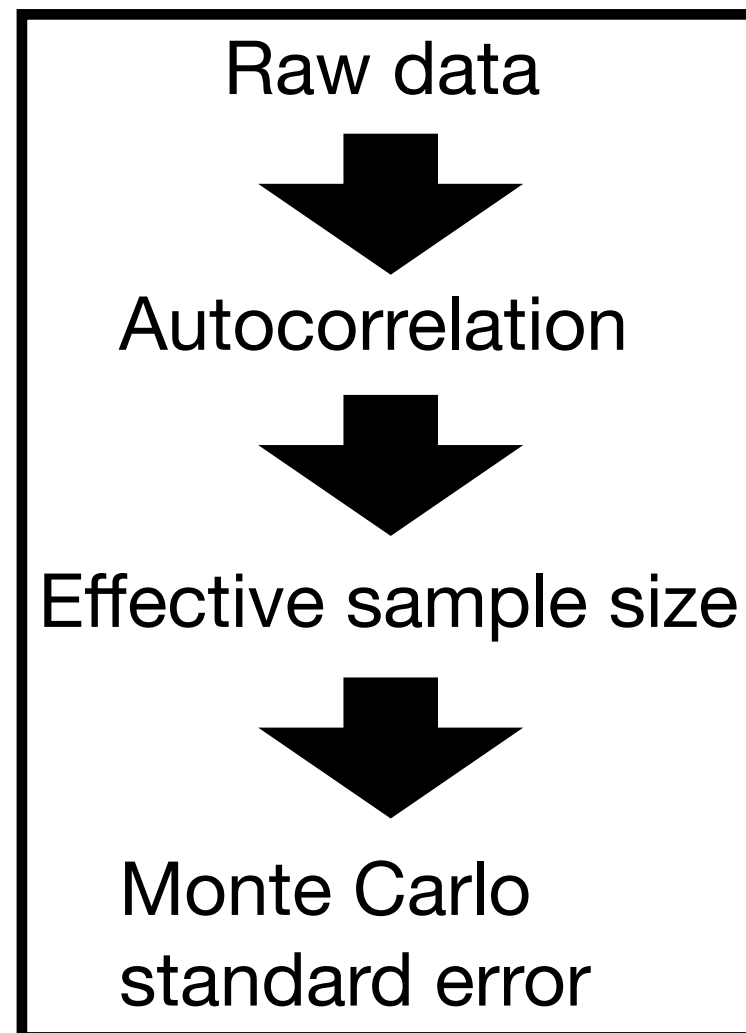
Similar to conventional definition

Performance evaluation

Sample more gives “Accuracy”



Quick review




Performance evaluation

Sample more gives “Accuracy”

The ‘accuracy’ also depends on which quantity wanted



With the same sample size



Performance evaluation

Get rid of initialization effects. “Representativeness”


Observe by eyes



A formal measurement

Sinking factor (Gelman-Rubin statistic)

$$\frac{\text{Variance between traces}}{\text{Variance within traces}} \rightarrow \text{towards 1}$$



Performance evaluation

Practical technics. “Efficiency”

Parallel hardware

Adjust the sampling method

Change the parameterization

MCMC Algorithms in Action

ANIMATION

<https://chi-feng.github.io/mcmc-demo/>

APPENDICES

How Do We Sample – A Quick Remark on U(0,1)

Algorithm: A (Lousy) RNG – Coin Flipping

- ① Flip a coin N times (e.g. 20) and record the results in sequence



1	0	0	0	0	0	0	1	1	0
1	0	0	0	0	1	1	0	0	1

Binary

(530969)

Decimal

- ② Divide the (decimal) result by the maximum number possible

Sampled Random Number: $x = \frac{530969}{1048575} = 0.506373$


All sampling is fundamentally constructed from a uniform RNG!

The Cumulative Density Function (CDF)

RECALL:

The cumulative density function is given by

$$F_X(\alpha) = \mathbb{P}(x \leq \alpha) = \int_{-\infty}^{\alpha} p_X(x) dx$$




The range of the CDF conveniently lies **between 0 and 1**

What is the **density of points** along the y-axis?

A First Attempt at Examining the y-axis Density

Let us begin by naively considering **evenly spaced points** along the x-axis...

This is called a “rug plot”, commonly used to visualize the distribution of the data



The points are clearly not evenly spaced along the y-axis


The spacing depends on how '**steep**' the curve is

QUESTION:

Is there something wrong here?

A Correct Examination of the y-axis Density

CORRECTION: Since this is a CDF, the density of points along the x-axis should depend on the **PDF of x** !



The points seem to be evenly spaced along the y-axis!

Can we prove this?
Can we use this result to do anything useful?

(HOMEWORK EXERCISE)

Inverse CDF Transform Sampling

IDEA: Starting with samples drawn from the PDF of x , we end up with uniformly distributed samples when applying the CDF to them.

Why not reverse the process?

Algorithm: Inverse CDF Sampling

①

Calculate the inverse of the target CDF F_X^{-1}

Not possible
in general

②

Sample iid random variables from a uniform distribution $y^i \sim \mathcal{U}(0,1)$


③

Transforming each sample via $x^i = F_X^{-1}(y^i)$


Inverse CDF Transform Sampling

(HOMEWORK EXERCISE)

Example: Sampling from the Laplace distribution via inverse CDF transform sampling



Rejection Sampling: Why Can We Do “2D Shooting”?



We can ‘smear out’ the samples evenly in the y-direction – how much we can ‘smear’ depends on how much density we had at that point

This ‘smearing out’ makes every 2D square have **equal probability density**

Mathematically, the PDF of x can be considered as the **marginal distribution** of x wrt to the joint PDF

$$p(x, z) = \begin{cases} 1 & \text{if } 0 \leq z \leq p_X(x) \\ 0 & \text{otherwise} \end{cases}$$

1	3	5	4	3	3	5	7	4	2
---	---	---	---	---	---	---	---	---	---